

RESEARCH INTERESTS

- **Summarization, Long-Context Generation, AI Safety, Model Efficiency, Reinforcement Learning**

EDUCATION

- **University of California, Riverside** Riverside, CA
Ph.D in Computer Science (Advisor: Yue Dong) *Sep. 2023 – Present*
- **Tianjin University** Tianjin, China
Master of Computer Technology (Advisor: Deyi Xiong) *Sep. 2020 – Jun. 2023*
- **Xidian University** Xi'an, China
Bachelor of Software Engineering *Sep. 2016 – June. 2020*

EXPERIENCE

- **Amazon** Palo Alto, CA
Applied Research Intern *Jun. 2025 - Sep. 2025*
 - **Project:** Mitigating Hallucination for KV-cache compression (One paper submitted to ACL2026)
- **Microsoft** Seattle, WA
Research Intern *Jun. 2024 - Sep. 2024*
 - **Project:** KV-cache Compression for Long-Context Generation (One paper accepted to ICLR2025)
- **University of California, Riverside** Riverside, CA
Research Intern *Sep. 2022 - Aug. 2023*
 - **Project:** Inverse Reinforcement Learning for Summarization. (One paper accepted to EMNLP2023)
- **LangBoat** Beijing, China
Research Intern *Jan. 2022 - Sep. 2022*
 - **Project:** Neural Machine Translation with Translation Memory

PUBLICATIONS

- [1]: **Yu Fu**, Haz Sameen Shahgir, Hui Liu, Xianfeng Tang, Qi He, Yue Dong (2026). Harnessing the Unseen: The Hidden Influence of Intrinsic Knowledge in Long-Context Language Models
AAAI 2026
- [2]: **Yu Fu**, Zefan Cai, Abedelkadir Asi, Wayne Xiong, Yue Dong, Wen Xiao (2025). Not All Heads Matter: A Head-Level KV Cache Compression Method with Integrated Retrieval and Reasoning.
ICLR 2025
- [3]: **Yu Fu**, Yufei Li, Wen Xiao, Cong Liu, Yue Dong (2024). Cross-Task Defense: Instruction-Tuning LLMs for Content Safety
NAACL 2024@TrustNLP
- [4]: **Yu Fu**, Yufei Li, Wen Xiao, Cong Liu, Yue Dong (2024). Safety Alignment in NLP Tasks: Weakly Aligned Summarization as an In-Context Attack.
ACL 2024
- [5]: Yiran Luo, Het Patel, **Yu Fu**, Dawon Ahn, Jia Chen, Yue Dong, Evangelos E. Papalexakis (2024). TRAWL: Tensor Reduced and Approximated Weights for Large Language Models.
PAKDD 2025@DSFA
- [6]: Erfan Shayegani, Md Abdullah Al Mamun, **Yu Fu**, Pedram Zaree, Yue Dong, Nael Abu-Ghazaleh. Survey of Vulnerabilities in Large Language Models Revealed by Adversarial Attacks.
ACL 2024 tutorial
- [7]: **Yu Fu**, Deyi Xiong, Yue Dong (2023). Watermarking Conditional Text Generation for AI Detection: Unveiling Challenges and a Semantic-Aware Watermark Remedy.
AAAI2024

- [8]: **Yu Fu**, Deyi Xiong, Yue Dong (2023). Inverse Reinforcement Learning for Text Summarization. **Findings of EMNLP2023**
- [9]: Jie He, **Yu Fu** (2022). MetaXCR: Reinforcement-Based Meta-Transfer Learning for Cross-Lingual Commonsense Reasoning. **PMLR@TLNP**
- [10]: Fan Deng, Zhenhua Yu, Wenjing Liu, Xiaoqing Luo, **Yu Fu**, Ben Qiang, Chaoyang Xu, Zhiwu Li (2021). An efficient policy evaluation engine for XACML policy management. **Information Sciences Journal**

HONORS & AWARDS

- **UCR Dean's Distinguished Fellowship** Riverside, CA
Personal *Sep. 2023 – Sep. 2024*

PROGRAMMING SKILLS

- **Programming:** Python, Pytorch, Faiseq, Transformers, RL for Summarization
- **Natural Languages:** Mandarin Chinese (Native), English (Proficient)

PROGRAM COMMITTEE & REVIEWER

- AAAI2026, ACL Rolling Review, EMNLP2024, NAACL-HLT/ACL/EMNLP Tutorials 2025